

BREEZE: Training-Free and Auditable LLM Recipe Augmentation with Physics-Guided Admission

Author One^a, Author Two^a

^aDepartment, City, Country

ARTICLE INFO

Keywords:

Condition monitoring
Large language model
Synthetic data
Physics-guided admission
Few-shot diagnosis
Time-series generation

ABSTRACT

When labelled target faults are scarce, fitting, selecting, and maintaining a target-specific waveform generator adds a separate optimization pipeline to few-shot diagnosis. We present BREEZE, a training-free, physics-guided synthetic fault-data generation framework in which an LLM selects structured recipes from supplied kinematics, training-fold statistics, and exemplar descriptions, and a fixed seeded renderer constructs the waveforms. A verifier calibrated only on the outer-training split admits candidates, returns bounded feedback for recipe revision, and archives recipe, seed, waveform-hash, and gate evidence, requiring neither target-specific generator optimization nor LLM fine-tuning. Class-conditional RMS, kurtosis, PSD, band-energy, and envelope-frequency diagnostics expose both fidelity gains and failures: on PU, the LLM pool reduces RMS, PSD, and band-energy errors relative to rule and random open-loop sources, while no source dominates every diagnostic. Across frozen, leakage-resistant protocols on PU, CWRU, and Berkeley, the admitted pools yield paired gains over repeated few-shot subset draws and CNN initializations: LLM recipes exceed rule and random open-loop sources on PU, rule, noise augmentation, and real-only in provenance-valid CWRU comparisons, and both non-structured Berkeley baselines at every shot, with only a 0.2–0.5-point advantage over rule in significant lower-shot cells and convergence at ten shots. Six predeclared PU cross-condition attempts delimit the current morphology-transfer boundary; together, the results position explicit, auditable, training-free generation as a practical few-shot augmentation operating point.

1. Introduction


Industrial fault diagnosis faces a structural data imbalance. Healthy operation is abundant, whereas labelled faults are scarce, expensive to seed, and rarely observed under every load and speed. New machines, sensor layouts, and operating regimes therefore offer only a few examples of the events that a diagnostic model must recognize.

Generator-based augmentation can alleviate this scarcity, but fitting a model for each target machine or condition introduces a separate optimization and model-selection problem. Under scarce or imbalanced faults, DFE-GAN addresses adversarial instability, mode collapse, and weak sample filtering through a modified objective, attention, and an automatic filter (Liu, Jiang, Wu and Li, 2022). Diffusion removes the adversarial game but still fits a target distribution: TSDM redesigns the original DDPM denoiser for one-dimensional vibration structure, ReF-DDPM strengthens intra- and inter-level features, a TII study adds a comprehensive evaluation framework, and FaultDiffusion attributes few-shot failure to the normal–fault domain gap and high intra-class variation (Yi, Hou, Jin, Saeed, Kandil and Duan, 2024; Yu, Li, Huang, Xiao, Zhang, Li and Fu, 2024; Yang, Ye, Yuan, Zhu, Mei and Zhou, 2024; Xu, Chen, Wang, Li, Tang, Zhao and Liu, 2025). ACS-DM and MR-DDIM further modify trained denoisers and objectives to balance temporal diversity with spectral fidelity (Tong, Zhu, Lu, Wang, Du, Jiang, Xu, Huo and Du, 2026; Xu, Tong, Zhu, Du, Zhao, Jiang and Wang, 2026).

These approaches can be effective when a representative target corpus supports distribution fitting, design choices, and checkpoint selection. In the few-shot regime, however, the same limited target evidence must serve all three roles. Even when impact priors or spectral constraints enter the objective, they influence generated signals through target-data optimization and model selection (Gao, Xu, Zhang, Fang and Tan, 2025; Liu, Han, Chen, Ding, Zhang, Shen, Wu, Zhang and Li, 2026). We study a different engineering operating point: remove target-specific generator fitting from augmentation and keep the controllable signal description explicit.

BREEZE is a synthetic fault-data generation framework that changes where controllable generation choices reside. First, the LLM selects kinematics, impacts, modulation, and spectral allocation in a structured JSON recipe using prompt-supplied machine information, training-fold statistics, and exemplar descriptions; a fixed seeded renderer

a

Closed-loop physics-verified admission

No target-data generator optimization and no waveform repair: rejected candidates remain rejected.

Figure 1: The BREEZE workflow. The LLM proposes a structured recipe, a fixed seeded renderer constructs vibration and, when available, current channels, and a verifier calibrated exclusively on the real training split admits or rejects the candidate. A rejection becomes structured feedback for at most K further rounds. Rejected signals are not repaired.

executes those fields as explicit signal equations. The evidence therefore tests prompt-supported, LLM-mediated recipe selection rather than whether pretraining alone contains fault physics. Second, a verifier moves physical control from a learned objective to gates calibrated on the corresponding outer-training split, and measured failures condition bounded recipe revision. Third, recipe, seed, waveform hash, and gate report records make every admitted sample traceable. The verifier is not a stand-alone post hoc audit: it closes the generation loop by determining which rendered candidates enter the augmentation pool and which recipes receive bounded revision. Figure 1 summarizes these components.

Three evaluation questions test this causal chain. First, can random legal values or a fixed engineer-written rule replace the LLM's recipe knowledge when renderer and verifier are held constant? Second, how does train-calibrated closed-loop admission shape feasible pools and their physical evidence? Third, do the resulting pools improve file-, load-, or experiment-disjoint few-shot diagnosis, and where does LLM-mediated recipe selection stop transferring? Paired source ablations, admission records, class-conditional diagnostics, leakage-resistant splits, and corrected tests answer the questions together.

The paper makes four contributions:

1. **Introduce** a training-free synthetic fault-data generation framework that combines prompt-supported LLM recipe generation, deterministic waveform rendering, and verifier-guided refinement without optimizing a target-data generator.
2. **Provide** explicit physical control and sample-level auditability through train-calibrated admission over robust statistics, Welch-PSD shape, envelope-spectrum evidence, reliability-gated current evidence, and diversity, with bounded feedback and archived evidence.
3. **Establish** complementary physical-fidelity and downstream-utility evidence while isolating the LLM-mediated recipe contribution through a verifier-matched rule source and a distinct random open-loop control on PU, CWRU, and Berkeley. The registered evidence comprises 20 PU seeds and 40 seeds for CWRU and Berkeley, with family-wise Holm correction and a global BH sensitivity audit.

4. **Delineate** the operating boundary through the complete six-stage PU leave-one-condition-out (LOCO) failure chain, process-metadata confounding in UMich milling, and a stopped MU-TCM inner-validation line.

The evidence supports a bounded generation claim: BREEZE produces physically structured synthetic fault signals that satisfy declared train-supported gates, and the resulting fixed pools improve the stated few-shot tasks under their frozen protocols. The scope is training-free augmentation, not training-free diagnosis; few-shot real calibration and a trained downstream classifier remain necessary.

2. Related work: where generative knowledge resides

2.1. Knowledge learned from target signals

General time-series generators place distributional knowledge in learned weights. Recurrent conditional GANs, TimeGAN, transformer GANs, variational models, and diffusion models all fit a target sequence corpus (Esteban, Hyland and Rätsch, 2017; Yoon, Jarrett and van der Schaar, 2019; Li, Metsis, Wang and Ngu, 2022a; Li, Ngu and Metsis, 2022b; Liao, Ni, Szpruch, Wiese, Sabate-Vidales and Xiao, 2020; Desai, Freeman, Wang and Beaver, 2021; Nichol and Dhariwal, 2021; Tashiro, Song, Song and Ermon, 2021; Alcaraz and Strodthoff, 2022; Rasul, Seward, Schuster and Vollgraf, 2021). Machinery-specific GANs improve this fitting through sparsity, feature enhancement, frequency learning, spectrum-guided sampling, and limited-data architectures (Ma, Ding, Wang, Wang, Ma and Lu, 2021; Liu et al., 2022; Ko, Lee, Kim, Kim and Youn, 2023; Kim, Ko, Lee, Kim, Jung and Youn, 2024; Lin, Kong, Zhao, Han, Dong, Liu and Chu, 2025). Recent diffusion studies replace adversarial learning with trained denoising while redesigning feature paths or sampling for vibration data (Yi et al., 2024; Yu et al., 2024; Yang et al., 2024; Xu et al., 2025; Tong et al., 2026). These methods retain knowledge in fitted generator weights; BREEZE instead evaluates an operating point without target-generator optimization.

2.2. Physical knowledge learned through objectives and architecture

Physics-informed generation adds domain knowledge to the learning process. Equations, mechanisms, or physical penalties constrain a trainable model (Raissi, Perdikaris and Karniadakis, 2019); rotating-machinery systems encode impact response, speed modulation, multimodal structure, condition embeddings, or spectral consistency (Gao et al., 2025; Liu et al., 2026; Guo, Xu, Zheng, Xie and Wang, 2025; Xu et al., 2026). The prior resides in an objective or architecture, but its effect on generated signals still passes through optimized weights. BREEZE places the signal equations in a fixed renderer and evaluates physical evidence after rendering through train-calibrated admission.

2.3. Description-conditioned learned generation

Language descriptions provide another route for conditioning a learned model. LLMs have been used as numerical time-series forecasters and supplied with machinery priors for bearing health management (Gruver, Finzi, Qiu and Wilson, 2023; Peng, Liu, Du, Gao and Wang, 2025). BearGen is the closest signal-generation study: an LLM produces a signal description that conditions a diffusion generator trained across multiple bearing datasets (Lee, Jeon, Kim and Kim, 2026). Its description enters generation through learned denoising weights; BREEZE asks for an executable numerical recipe and uses a fixed renderer whose target-data path contains no learned generator weights.

2.4. Explicit knowledge with calibrated admission

Generated-data filtering controls which fitted samples reach a classifier. DFE-GAN and SGAN-DDS show that filtering, spectrum guidance, and diversity sampling affect the retained pool (Liu et al., 2022; Kim et al., 2024). BREEZE combines an explicit recipe, fixed equations, train-only gate calibration, structured feedback, and per-sample records. Physical knowledge therefore resides in the recipe and renderer, while training data determine the supported admission region rather than generator weights.

Few-shot transfer and meta-learning address a complementary question: how a diagnostic classifier acquires or adapts representation knowledge from related tasks and source domains. Bearing and machinery studies use episodic or meta-trained adaptation under limited target labels and varying conditions (Li, Li, Zhang, He, Liao and Hu, 2021; Wu, Zhao, Sun, Yan and Chen, 2020; Su, Xiang, Hu, Xu and Yang, 2022; Feng, Chen, Xie, Zhang, Lv and Pan, 2022). BREEZE does not replace those classifier-side methods; it tests whether an auditable synthetic pool can improve the stated few-shot protocol without fitting a target-data generator.

Table 1

Positioning by where generative and physical knowledge resides. “Target training” denotes optimization of a generator on the target signal corpus.

Method family	Generative mechanism	Target training	Where physical knowledge resides	Physical control	Primary evidence pattern
DFE-GAN / SGAN-DDS	Learned adversarial generator plus filtering or sampling	Yes	Generator and discriminator weights	Learned features, filtering, and spectrum guidance	Fidelity, diversity, and diagnosis
Physics-informed GAN / diffusion	Learned generator with physical priors or condition embedding	Yes	Physical objective, architecture, and optimized weights	Training loss or architecture	Mechanism, multidomain fidelity, diagnosis
BearGen	LLM descriptions condition a learned diffusion model	Yes	Description prompt and learned denoising weights	Description-conditioned denoising	Eight bearing datasets, quality, few-shot, cost
BREEZE	LLM JSON recipe plus fixed seeded renderer	No	Recipe, renderer equations, and calibrated gate thresholds	Train-calibrated admission and rejection	Source ablation, gate reports, physical distances, paired diagnosis

2.5. Diagnostic observables used as gates

Established diagnostic observables expose complementary aspects of rolling-element faults. Envelope analysis, spectral kurtosis, and the fast kurtogram reveal stochastic or pseudo-cyclostationary impacts (Randall and Antoni, 2011; Antoni, 2006, 2007), while Welch PSD estimation represents spectral shape (Welch, 1967). Motor-current sidebands add bearing evidence when the real training data support class separability (Schoen, Habetler, Kamran and Bartheld, 1995). BREEZE differs by converting these observables into explicit, train-calibrated admission gates.

3. Problem formulation and framework

3.1. Train-calibrated admission

The second component moves physical control out of generator training and into a decision rule calibrated entirely within the outer training split. Let $\mathcal{D}_{\text{tr}} = \{(x_i, y_i, m_i)\}_{i=1}^n$ denote its real windows, class labels, and available machine metadata. For requested class y and condition m , the LLM proposes recipe r_t at feedback round t . The following mapping makes recipe execution and local randomness explicit:

$$\tilde{x}_{t,s} = R(r_t, m; s), \quad (1)$$

where the seeded renderer R archives s with the recipe. The verifier then maps the rendered waveform to both an admission decision and its evidence report:


$$(a_{t,s}, q_{t,s}) = V(\tilde{x}_{t,s}, y, m; \mathcal{D}_{\text{tr}}), \quad a_{t,s} \in \{0, 1\}. \quad (2)$$

Equation (2) exposes the decision $a_{t,s}$ and the diagnostic report $q_{t,s}$ as separate outputs. Renderer inputs, gate calibration, feedback, and pool selection all remain confined to the outer training split.

Admission must represent interval support, direction-specific evidence, healthy suppression, and diversity in one decision. Partition the active feature gates into two-sided, upper-bound, and lower-bound index sets $\mathcal{J}_{\text{int}}$, $\mathcal{J}_{\text{up}}$, and $\mathcal{J}_{\text{low}}$. For deterministic feature f_j , class-supported bounds $[l_{j,y,m}, u_{j,y,m}]$, healthy fault-prominence score p_f , physical embedding $e(x)$, and current admitted set $\mathcal{S}_y$, the implemented predicate types are

$$\begin{aligned} l_{j,y,m} &\leq f_j(x) \leq u_{j,y,m}, & j \in \mathcal{J}_{\text{int}} & \text{ (interval),} \\ f_j(x) &\leq u_{j,y,m}, & j \in \mathcal{J}_{\text{up}} & \text{ (upper-bound evidence),} \\ f_j(x) &\geq l_{j,y,m}, & j \in \mathcal{J}_{\text{low}} & \text{ (lower-bound evidence),} \\ p_f(x) &\leq u_{f,\text{healthy}} & & \text{ (healthy suppression),} \end{aligned}$$

a Recipe-renderer-verifier responsibility boundary



Only the highlighted component owns each output; evidence is not transferred across responsibility boundaries.

Figure 2: Component responsibility boundary. The LLM selects recipe parameters; the renderer alone constructs waveforms; the verifier admits or rejects without repair; and the classifier is used only for downstream evaluation. Every formal source comparison holds renderer and classifier fixed.

$$S_y = \emptyset \text{ or } \min_{z \in S_y} \|e(x) - e(z)\|_2 \geq \delta_y \quad (\text{diversity}). \quad (3)$$

The empty-set clause admits the first candidate with respect to diversity only; all active per-candidate gates must still pass. Subsequent candidates must meet the distance lower bound. The intersection of the active predicates defines the admissible set. Membership therefore records compliance with the configured gates under their training calibration.

3.2. Responsibility boundaries

A component claim is meaningful only when proposal knowledge, waveform construction, admission, and diagnosis have explicit responsibility boundaries. Figure 2 assigns one component to each responsibility. The LLM selects recipe fields, the renderer enforces the signal equations, the verifier measures compatibility with train-supported evidence, and the downstream classifier measures utility. Controlled source comparisons can then isolate the recipe source while holding the other components fixed.

The structured recipe exposes every LLM-controlled degree of freedom. It records shaft and fault frequencies, target RMS, periodic-impact rate, amplitude, decay, resonance, slip jitter, amplitude variation, load-zone modulation, background components, random machine transients, eight vibration-band weights, and optional two-phase current parameters. The initial prompt supplies class, condition metadata, bearing kinematics, a training-derived exemplar description, and the JSON schema. The feedback prompt adds the previous recipe, each measured gate violation, its supported bound, and a request for proportional parameter change. Complete prompt builders and schemas are preserved in `src/llm.py`.

3.3. Seeded renderer

The first component can bypass learned generator weights only when every recipe field enters the waveform through fixed equations. For PU vibration, the renderer combines deterministic background harmonics, Poisson machine transients, a jittered train of resonant impacts, and band-shaped noise. The equations expose the source terms and the link between impact interval and requested fault frequency:

$$x_v(t) = \sum_q A_q \sin(2\pi f_q t + \phi_q) + \sum_{\ell} b_{\ell} h(t - u_{\ell}) + \sum_k a_k h(t - t_k) + \epsilon_b(t), \quad (4)$$

$$h(\tau) = \mathbf{1}_{\tau \geq 0} \exp(-\tau/\tau_d) \sin(2\pi f_r \tau + \phi), \quad (5)$$

$$t_{k+1} - t_k = f_f^{-1}(1 + \xi_k). \quad (6)$$

Here u_ℓ follows a Poisson background process, f_f is the requested fault frequency, and ξ_k is bounded recipe jitter. Inner-race impact amplitudes may be modulated at shaft frequency. The renderer shapes additive noise over the fixed eight-band grid and applies the same deterministic target- RMS normalization to every recipe source before admission. Both operations belong to R and precede the gate decision.

When two phase currents are available, the following equation makes carrier, fault modulation, harmonics, and channel noise explicit for channel p :

$$i_p(t) = A_p [1 + \beta \cos(2\pi f_f t + \psi)] \left[\sin(2\pi f_e t + \phi_p) + \sum_{h \in \{3,5,7\}} \rho_h \sin(2\pi h f_e t + \phi_{p,h}) \right] + \eta_p(t), \quad (7)$$

where f_e is supply frequency. The code and archived seed make $R(r, m; s)$ deterministic.

Berkeley requires a dataset-specific milling renderer under the same component responsibility principle. Its multichannel renderer starts from a train-only channel exemplar, adds the recipe's transient and band-energy structure, and preserves channel-specific location and scale. This renderer is calibrated inside the outer training fold, while Eqs. (4)–(7) define the bearing-channel model.

3.4. Verifier gates

The second component moves physical constraints from a generator-training loss to calibrated gates, so the verifier evaluates six complementary evidence families:

1. **Legality:** exact shape, finite samples, non-degenerate channels, and absence of long repeated runs.
2. **Robust statistics:** class- and channel-conditional RMS, standard deviation, peak, kurtosis, skewness, and crest factor inside train-derived quantile or robust-union domains.
3. **Spectral shape:** triangular-band Welch-PSD fractions and PSD-CDF Wasserstein-1 distance to the class training reference.
4. **Envelope evidence:** resonance-band Hilbert demodulation, prominence near the requested BPFO, BPFI, BSF, or FTF frequency, harmonic support, and suppression of unsupported peaks for healthy candidates.
5. **Current evidence:** vector-current sideband scores become active when training data establish class separability; otherwise they remain audit-only measurements.
6. **Diversity:** a lower nearest-neighbour bound in the physical feature embedding relative to a train-derived real–real reference.

PU training data leave the MCSA score audit-only. Current-signature consistency is therefore reported only for pools whose training split activates that gate.

3.5. Closed-loop construction and audit records

Measured failures shape the pool only when admission closes the loop around the proposal source. This loop must bound proposal cost while preserving every decision needed for audit and replay. Algorithm 1 states the contract for one proposal slot.

A *slot* is one requested class-conditioned recipe unit, whereas a *window* is the $C \times N$ tensor supplied to the classifier. One slot can archive candidates from several feedback rounds, and its admitted recipe can be rendered with several deterministic seeds. The frozen pipeline renders three seeds per recipe and can retain additional diversity-admitted expansions, so one admitted slot can yield multiple windows. This distinction matters because slot counts measure LLM proposal cost, while window counts measure the classifier's augmentation budget.

Generation provenance determines which stages can be replayed exactly. The CWRU archive records the OpenAI-compatible model identifier `mimo-v2.5`, temperature 0.8, top- p 0.9, and maximum 900 output tokens, although the provider did not expose sampling seeds. The earlier PU archive preserves prompts, recipes, rendered pools, and local seeds; its original provenance record omits the exact provider version and top- p . Consequently, the frozen PU pools are reproducible from archived recipes and renderer seeds, while exact regeneration of the original language responses remains unavailable. Manuscript assembly required no new LLM call.

Algorithm 1 BREEZE recipe admission for one proposal slot**Require:** class y , condition m , LLM G , renderer R , verifier V , round limit K , renderer seeds S

```

1:  $q_{-1} \leftarrow \emptyset$ 
2: for  $t = 0$  to  $K$  do
3:    $r_t \leftarrow G(y, m, q_{t-1})$ 
4:   for all  $s \in S$  do
5:      $\tilde{x}_{t,s} \leftarrow R(r_t, m; s)$ 
6:      $(a_{t,s}, q_{t,s}) \leftarrow V(\tilde{x}_{t,s}, y, m)$ 
7:     if  $a_{t,s} = 1$  then
8:       archive recipe, seed, waveform hash, and full gate report
9:     return admitted slot and all diversity-admitted windows
10:   end if
11: end for
12:  $q_t \leftarrow \text{AggregateFailures}(\{q_{t,s}\})$ 
13: end for
14: return rejected slot with terminal failure report

```

Table 2

Formal datasets and frozen protocols. Counts are windows after the declared split.

Dataset	Sampling and input	Classes	Operating regime	Train/test	Split unit
PU	8 kHz, 3 channels, 2048 samples	healthy, OR, IR	N09_M07_F10: 900 rpm, 0.7 Nm, 1000 N	3846/962	chronological measurement files within bearing; zero shared file/window keys
CWRU	12 kHz drive-end vibration, 2048 samples	healthy, IR, ball, OR	motor loads 0–3; defects 0.007–0.028 in	696/306 for within-load0	chronological within-file for load0; loads 1–3 held out separately for clean transfer tests
Berkeley	250 Hz, 6 channels, 512 samples	healthy ($VB < 0.2$), degraded ($VB \geq 0.2$)	feed, depth, and material combinations	1836/646	disjoint machining case/run groups

4. Experimental design

4.1. Public datasets and leakage-resistant splits

Leakage-resistant evaluation requires splitting acquisition units before sampling windows. Window-level random splits can mix recording-specific or overlapping signal structure across partitions and thereby answer a weaker generalization question; a recent bearing analysis documents materially split-dependent conclusions, while a broader review classifies the resulting evaluation risks (Wheat, Mohrenschildt, Habibi and Al-Ani, 2024; Kapoor and Narayanan, 2023). The formal evidence uses the established PU and CWRU bearing benchmarks (Lessmeier, Kimotho, Zimmer and Sextro, 2016; Smith and Randall, 2015) and Berkeley milling from the NASA prognostics repository (Agogino and Goebel, 2007). Their file-, load-, or case/run-level protocols appear in Table 2, and only these grouped splits support formal claims.

The PU protocol isolates measurement files within each bearing. Its healthy, OR, and IR classes contain 1200, 1202, and 1444 training windows, together with 300, 301, and 361 test windows. The input retains one vibration and two phase-current channels.

The CWRU protocol tests both within-load diagnosis and source-load transfer. It evaluates a within-load0 split and three folds that hold out load1, load2, or load3. The frozen synthetic pools derive from load0, which belongs to the training side of these three folds. Reusing these pools gives the archived held-out-load0 fold target-load provenance, so that fold remains an audit artifact.

The Berkeley protocol isolates machining case/run groups before windowing. Its six channels contain spindle AC and DC current, table and spindle vibration, and table and spindle acoustic emission. Train-only inner validation fixed the healthy/degraded threshold and split before the one-time 40-seed formal run.

4.2. Recipe-source and augmentation controls

Testing whether LLM-mediated recipe selection adds value without target-generator learning requires changing the proposal source while holding every downstream factor fixed. The central ablation therefore retains the applicable renderer, data split, synthetic budget, CNN, and paired seeds:

- **LLM + verifier:** class- and condition-aware JSON recipe, physical admission, and bounded feedback.
- **Rule + verifier:** engineer-written class template with train-supported parameter ranges, the same renderer, and the same verifier; no LLM call.
- **Random open-loop:** random legal recipe values rendered without feedback. On PU this is an open-loop non-semantic downstream control; a separate random-recipe verifier audit is also reported.
- **Noise augmentation:** real few-shot windows with channel-scaled jitter and amplitude scaling, using the same synthetic count.
- **Real only:** the paired few-shot subset without synthetic data.

The completed formal families match these controls to each dataset. CWRU compares LLM with rule, noise, and real only over within-load0 and held-out loads 1–3. Berkeley adds random open-loop, while PU Phase-A focuses on LLM versus rule and random open-loop.

4.3. Few-shot classifier and statistical protocol

Matched downstream evaluation requires a fixed classifier and paired sources of randomness. All methods within a dataset therefore share the same compact 1-D CNN. Its three convolution blocks use 16, 32, and 64 channels, kernels 15, 9, and 5, stride two, max-pooling after the first two blocks, adaptive average pooling to length eight, and a linear classifier. Adam uses learning rate 3×10^{-4} , weight decay 10^{-4} , batch size 32, and cross-entropy loss. PU uses 60 epochs; the frozen CWRU and Berkeley protocols use 20. Training-channel mean and standard deviation normalize both train and test data.

The dataset-specific budgets define the registered comparison cells. PU uses $n_{\text{real}} \in \{5, 10, 25\}$, 150 synthetic windows per class, and 20 paired seeds. CWRU uses the same shot counts, 20 synthetic windows per class, and 40 seeds. Berkeley uses $n_{\text{real}} \in \{2, 5, 10\}$, 20 synthetic windows per class, and 40 seeds. Accuracy and Macro-F1 are primary, while each paired seed selects the same real subset and classifier initialization for every method in its cell.

The statistical unit is a paired downstream seed, combining a CNN initialization with a few-shot subset draw. Within each dataset protocol, every seed reuses the same fixed synthetic pool for a method; seeds are therefore not independent LLM generations or independently generated pools. The resulting inference is precisely that the pool produced by the archived generation process is stably useful across repeated few-shot selection and classifier training. It does not estimate variation across independently generated pools.

Registered inference preserves the pairing used in training. One-sided Wilcoxon signed-rank tests compare paired seeds (Wilcoxon, 1945), and Holm correction controls family-wise error within each predeclared comparison family (Holm, 1979). The registered decision follows Holm. As a cross-dataset sensitivity analysis, we additionally apply one global Benjamini–Hochberg correction to the 102 core PU, provenance-valid CWRU, and Berkeley hypotheses (Benjamini and Hochberg, 1995). Table 3 shows that this global FDR analysis preserves all 102 registered pass/fail decisions; it does not replace the preregistered family structure. We report direction, adjusted decision, seed count, and the complete pass pattern. An unadjusted p value near 0.05 remains descriptive; significance requires the adjusted Holm decision.

4.4. Physical and distribution diagnostics

Physical and distribution diagnostics complement downstream accuracy by measuring how each class relates to its outer-training reference. They include RMS, kurtosis, skewness, crest factor, peak value, BPFO/BPFI/BSF/FTF peak alignment where applicable, envelope prominence and harmonic support, PSD-CDF Wasserstein-1 distance, relative

Table 3

Multiplicity sensitivity over the 102 core paired hypotheses. Registered within-family Holm decisions remain primary; global BH applies one FDR family across PU, provenance-valid CWRU, and Berkeley.

dataset	hypotheses	Holm pass	global BH pass	agreement
PU	12	12	12	12/12
CWRU	72	72	72	72/72
Berkeley	18	15	15	18/18
total	102	99	99	102/102

Table 4

PU Phase-A v2 file-split results (20 seeds; $n_{\text{syn}} = 150$ per class). Holm q values test LLM $K = 3$ superiority within each registered family.

n_{real}	metric	LLM $K = 3$	rule	random	q_{rule}	q_{random}
5	Acc.	0.777 ± 0.027	0.696 ± 0.053	0.450 ± 0.029	<0.001	<0.001
5	Macro-F1	0.776 ± 0.026	0.691 ± 0.049	0.436 ± 0.053	<0.001	<0.001
10	Acc.	0.796 ± 0.030	0.724 ± 0.042	0.483 ± 0.013	<0.001	<0.001
10	Macro-F1	0.800 ± 0.029	0.724 ± 0.040	0.488 ± 0.015	<0.001	<0.001
25	Acc.	0.818 ± 0.021	0.769 ± 0.025	0.529 ± 0.022	<0.001	<0.001
25	Macro-F1	0.821 ± 0.022	0.770 ± 0.024	0.527 ± 0.017	<0.001	<0.001

band-energy error, and nearest-neighbour distance. Wasserstein distance follows the optimal-transport interpretation of Villani (2009). The archived feature audit retains MMD with an RBF kernel and median pairwise-distance bandwidth (Gretton, Borgwardt, Rasch, Schölkopf and Smola, 2012), while the main figure reports directly interpretable distances. The core evidence uses these direct diagnostics; t-SNE, UMAP, and Mixup remain outside the frozen formal comparison.

To isolate the verifier after recipe generation, a zero-API cached-candidate ablation re-screens the archived PU $K = 3$ proposals with M2 (statistics), M3 (spectral shape and PSD-W1), M4 (envelope evidence), or M5 (pool diversity) disabled one at a time. The full replay must reproduce the frozen raw and balanced pools before any variant is evaluated. Every evaluable variant retains the same $B = 150$ synthetic windows/class, file split, downstream CNN, and 20 paired seeds; a diversity threshold that cannot fill that fixed budget is reported as a capacity stop rather than resampled at a lower budget.

5. Results

The results follow the component chain established in the Introduction. PU isolates the recipe source, CWRU and Berkeley test complete generated pools, and the physical and admission audits examine how the verifier shapes those pools.

5.1. LLM-mediated recipes outperform rule and open-loop random sources on PU

The PU component ablation shows that LLM recipes outperform rule and random open-loop recipes at every registered shot count for both diagnostic metrics. All 12 pairwise tests pass Holm correction ($q < 0.001$), and Table 4 reports the paired protocol means and standard deviations.

Seed-level deltas confirm that the PU LLM advantage is positive across every registered source comparison for Accuracy and Macro-F1. Figure 3 reports the paired mean effects with bootstrap intervals across the three protocols. Because the renderer and verifier remain fixed, the comparison attributes the supported PU contribution to LLM-mediated recipe selection under the supplied prompt context within the complete BREEZE pipeline.

5.2. Generated LLM pools improve few-shot diagnosis within load0 and on held-out loads 1–3

The complete load0-derived LLM pools outperform every registered comparator within load0 and on provenance-valid held-out loads 1–3. The within-load0 means exceed rule and noise augmentation at each shown shot count




Figure 3: Paired downstream effects. Each marker is the paired mean LLM-comparator effect in percentage points over the registered seeds; lines are 95% percentile bootstrap intervals over paired seeds. (a–c) LLM minus rule on PU, CWRU within-load0, and Berkeley under one shared ± 12 -point scale. (d–f) LLM minus random open-loop on PU and LLM minus noise augmentation on CWRU and Berkeley, with per-panel ranges printed. Circles are Accuracy and squares Macro-F1; filled markers pass the registered Holm correction and hollow markers do not. Seeds repeat few-shot subset selection and CNN initialization around one fixed pool per source.

Table 5

CWRU within-load0 results (40 seeds). All 72/72 provenance-valid LLM-comparator tests across within-load0 and held-out loads 1–3 pass Holm correction. The archived held-out-load0 fold is excluded because it reuses a load0-derived synthetic pool.

n_{real}	metric	LLM	rule	noise augmentation
5	Acc.	0.564 $\pm$ 0.068	0.549 $\pm$ 0.053	0.454 $\pm$ 0.066
5	Macro-F1	0.567 $\pm$ 0.075	0.554 $\pm$ 0.062	0.460 $\pm$ 0.084
10	Acc.	0.621 $\pm$ 0.037	0.584 $\pm$ 0.033	0.534 $\pm$ 0.065
10	Macro-F1	0.646 $\pm$ 0.050	0.611 $\pm$ 0.046	0.554 $\pm$ 0.082
25	Acc.	0.674 $\pm$ 0.037	0.652 $\pm$ 0.037	0.628 $\pm$ 0.055
25	Macro-F1	0.711 $\pm$ 0.030	0.689 $\pm$ 0.035	0.655 $\pm$ 0.053

(Table 5), while the combined protocol yields 72/72 positive Holm-corrected LLM-comparator decisions. Figure 9 reports the source-load0 LLM-rule transfer deltas: Accuracy improves by 1.7 to 3.5 percentage points, and Macro-F1 by 1.9 to 3.6 points across shots and held-out loads 1–3.

This result establishes source-load0 cross-load diagnostic utility for the tested held-out loads 1–3 and shows where explicit recipe knowledge remains useful beyond its source load under the supplied LLM context. A strict four-fold LOLO and cross-bearing scope remain outside this result because the reused synthetic pool gives the archived held-out-load0 fold target-load provenance.

Table 6

Berkeley binary milling formal test (40 seeds; $n_{\text{syn}} = 20$ per class). This is a partial result: all comparisons against noise augmentation and random open-loop pass, whereas LLM versus rule passes only Acc. at $n_{\text{real}} = 2$ and both metrics at $n_{\text{real}} = 5$.

n_{real}	metric	LLM	rule	noise aug.	random open-loop
2	Acc.	0.786 ± 0.021	0.783 ± 0.022	0.695 ± 0.110	0.524 ± 0.147
2	Macro-F1	0.778 ± 0.018	0.776 ± 0.019	0.675 ± 0.112	0.450 ± 0.127
5	Acc.	0.784 ± 0.030	0.779 ± 0.031	0.724 ± 0.068	0.614 ± 0.097
5	Macro-F1	0.773 ± 0.031	0.769 ± 0.032	0.709 ± 0.071	0.514 ± 0.106
10	Acc.	0.787 ± 0.024	0.786 ± 0.024	0.750 ± 0.049	0.654 ± 0.062
10	Macro-F1	0.777 ± 0.024	0.776 ± 0.023	0.734 ± 0.056	0.572 ± 0.081

5.3. Generated LLM pools beat non-structured Berkeley controls and converge with rules at ten shots

The complete LLM pools beat the non-structured Berkeley controls in all 12 registered comparisons. Against the strong structured rule reference, Accuracy passes at $n_{\text{real}} = 2$, while the two-shot Macro-F1 comparison remains outside the Holm decision; both metrics pass at $n_{\text{real}} = 5$. At $n_{\text{real}} = 10$, the LLM and rule results converge and both registered comparisons fall outside the Holm decision. This complete pattern yields 15 of 18 passes, with means reported in Table 6.

The Berkeley pattern localizes the additional LLM recipe contribution to lower-shot settings, but the effect over rule is small: the passing mean deltas are 0.267 percentage points for two-shot Accuracy and 0.468/0.407 points for five-shot Accuracy/Macro-F1. At ten shots, the deltas contract to 0.055/0.020 points and do not pass Holm correction. The structured rule is therefore a strong reference, and convergence shows that structured recipe design explains most of the gain as real samples increase. This result supports only the qualified binary protocol.

5.4. Admitted PU pools preserve complementary evidence under pool-level diversity control

The verifier admits PU LLM pools that combine class-relevant waveform evidence with within-pool nearest-neighbour selection. Under the same fixed demodulation procedure, Fig. 4 compares medoid-selected real, LLM, rule, and random windows for the two fault classes together with population envelope spectra. Admitted LLM and rule recipes preserve impulsive and envelope structure more consistently than the random control while retaining visible stochastic variation. These traces illustrate the population-level distributions in Figs. 5–6.

The PU fidelity advantage of the LLM pool is concentrated in three key observables. The class-averaged RMS-W1, PSD-W1, and band-energy error in Table 7 are all lower for LLM than for both rule and random open-loop sources, indicating closer amplitude, spectral-distribution, and energy-allocation agreement under the frozen diagnostics. This advantage is not universal across physical evidence: noise augmentation has the lowest class-averaged RMS, kurtosis, and band-energy distances, consistent with its direct perturbation of real windows, and rule has the lower class-averaged fault-frequency alignment error. Per-class cells also include comparator-favouring PSD, kurtosis, and alignment results. Table 8 therefore retains every bearing kurtosis/alignment cell rather than reducing the audit to a single realism ranking. The LLM nearest-neighbour value measures spacing among synthetic samples after scaling by real-reference feature variability. It is not a synthetic-to-real distance and therefore does not establish absence of training-window copying. This diversity value requires joint interpretation with fidelity because arbitrarily large distances can indicate unrealistic samples.

The separate synthetic-to-real audit in Table 9 uses the same physics-v3 pool selections and class-matched outer-training references. No available audited pool contains a byte-identical real training window, but several pools retain high maximum cross-correlation. Because no copying threshold was declared, the distances and correlations are reported as descriptive similarity diagnostics rather than proof of independence.

Berkeley uses a real train-only channel exemplar as the renderer background. Its rows therefore characterize how much carrier structure remains after the recipe-specific transient and band-energy additions; they do not support an independent-signal-generation claim.

These diagnostics test the admission component rather than rank generators in isolation. External downstream tests, controlled source comparisons, train/test separation, and failed cross-condition experiments add evidence beyond shared physical features, while gate compliance remains insufficient for physical truth.

Training-free LLM recipe augmentation




Figure 4: PU waveform and envelope-spectrum comparison for (a) outer-race and (b) inner-race classes; columns compare real training, LLM, rule, and random open-loop windows. Each waveform is the window nearest its source/class robust centre in a shared class-matched real-train-scaled feature space, with raw amplitude and shared class-specific y limits. Envelope spectra show the median and interquartile range over all windows of each source/class under the same fixed 500–2000 Hz demodulation band. Dashed markers denote the class-relevant kinematic frequency (BPFO 46.1 Hz; BPFI 73.9 Hz).

5.5. Closed-loop admission shapes feasible pools and exposes failure mechanisms

The closed-loop component separates proposal cost from the downstream augmentation budget while determining which recipes can form a feasible pool. PU requests 150 proposal slots per class; the round-level audit hashes all 450 archived slot records and assigns each admitted slot to its first candidate with `feasible=true`. It reproduces the frozen slot summary exactly. Figure 7(a) shows that 205 slots pass without feedback; rounds 1–3 add 36, 27, and 18 slots, respectively, reaching 286/450 by $K = 3$. At that point, the LLM process admits 60%, 60%, and 71% of healthy/OR/IR slots, while 164 slots remain unadmitted. These are complete slot counts for one frozen pool-generation run, not uncertainty estimates over independent LLM pools.

Recipe source determines whether the verifier can construct a balanced pool. The random-recipe audit admits no slot, whereas rule admission varies by class. This contrast makes the fixed rule a meaningful but incomplete substitute for LLM proposal under the PU verifier.

Terminal failures reveal class- and source-specific rejection mechanisms. Figure 8 reports the non-exclusive rates. Random recipes fail statistical and spectral gates much more often than cached LLM recipes. Healthy LLM candidates are rejected most often for unsupported envelope evidence, whereas fault candidates more often violate the robust




Figure 5: Class-conditional RMS (a–c) and kurtosis (d–f) empirical cumulative distributions for the real PU outer-training split and frozen synthetic pools. Real curves use every class-matched outer-training window (healthy/OR/IR: 1200/1202/1444), exactly the reference population used for the W1 diagnostics; synthetic curves use the fixed 150-window/class pools and are distinguished by both colour and line style. Admission narrows unsupported statistical excursions without forcing equality to a single reference value.

Table 7

PU physical diagnostics for fixed 150-window/class pools. RMS, kurtosis, PSD, and band-energy entries average all three classes; alignment averages the OR and IR fault classes because it is undefined for healthy windows. Lower is better for the five reference-relative diagnostics; diversity has no preferred direction.

pool	RMS W1	kurtosis W1	PSD W1	band-energy error	alignment [Hz]	NN diversity
llm	0.008	2.658	236.9	0.044	0.929	17.10
rule	0.031	3.425	271.6	0.235	0.547	14.20
random_open_loop	0.024	2.342	516.1	0.749	0.974	52.50
noise_aug	0.004	0.274	309.1	0.014	2.544	46.09

statistical domain. A slot may fail several gates, so the percentages are non-additive. The same report supplies the next feedback round with each measured violation and its bound. Closed-loop admission therefore turns calibrated evidence into an auditable pool-construction decision rather than a post hoc quality score.

5.6. Verifier components support the frozen PU pool rather than only rejecting candidates

The cached replay directly separates the gates from the LLM proposal source. Disabling M2, M3, or M4 expands slot admission, while disabling M5 leaves the per-slot decision unchanged and alters only pool-level diversity selection. Each single-gate removal produces negative paired Accuracy and Macro-F1 differences from the complete gate at every evaluated shot count (Table 10). M3 removal worsens the PSD diagnostic, whereas M4 removal can lower that one distance while still reducing downstream utility; thus no individual physical metric is treated as a universal quality score.

The $0.5\delta_y$ pool is byte-identical to the no-M5 pool under the fixed selection seed, so its downstream rows are an exact equivalence rather than a second independent training run. The reference threshold reproduces the frozen full pool. At $2.0\delta_y$, the cached healthy candidates cannot form the predeclared balanced budget, and no reduced-budget comparison is substituted.




Figure 6: Class-specific physical diagnostics on (a) PU, (b) CWRU, and (c) Berkeley. Each matrix cell is $\log_2$ of the method-to-rule distance ratio for RMS-W1, kurtosis-W1, PSD-W1, relative band-energy error, and fault-frequency alignment; negative values favour the method and positive values favour rule, on one shared zero-centred scale. NA denotes a frozen metric that is undefined or unavailable, not an imputed value. (d) Within-pool nearest-neighbour diversity in raw units on a separate log axis; diversity has no universally optimal direction and does not share the fidelity colour scale.

6. Cross-condition evidence and negative results

The root-cause framework also predicts where BREEZE should fail. Explicit recipes remove target distribution fitting from generator weights, but they do not identify an unknown mapping from operating condition to fault morphology. The following stress tests probe that missing knowledge, and their protocol value is to prevent extrapolation from positive within-condition results.

6.1. CWRU source-load transfer succeeds while PU morphology does not transfer

CWRU source-load0 recipes transfer to the tested loads 1–3, whereas the PU LOCO protocol remains unsuccessful. Figure 9 places these outcomes side by side: PU v1 fails 57/96 registered cells, and v2 fails 60/96. The PU failures cluster by held-out condition and persist across Accuracy, Macro-F1, and shot counts, locating the unresolved boundary in condition-dependent morphology.

The six-stage PU audit localizes the missing morphology–condition map from kinematics through source-only evidence (Fig. 10). v1 exposed a kinematic mismatch, and v2 corrected frequency handling while downstream

Table 8

Complete per-class bearing audit for kurtosis-W1 and fault-frequency alignment. Lower is better. Alignment is undefined for healthy windows (—). CWRU random open-loop is NA because no matching frozen physics-v3 pool exists; no value is imputed. Equal displayed values are retained rather than selectively suppressed.

dataset	pool	class	kurtosis W1	alignment [Hz]	dataset	pool	class	kurtosis W1	alignment [Hz]
PU	llm	healthy	4.495	—	CWRU	llm	B	3.881	0.466
PU	llm	OR	0.857	1.090	CWRU	llm	OR	7.865	2.055
PU	llm	IR	2.622	0.769	CWRU	rule	healthy	0.036	—
PU	rule	healthy	4.448	—	CWRU	rule	IR	5.929	2.064
PU	rule	OR	2.190	0.816	CWRU	rule	B	4.208	0.466
PU	rule	IR	3.637	0.277	CWRU	rule	OR	6.694	2.055
PU	random_open_loop	healthy	4.816	—	CWRU	noise_aug	healthy	0.043	—
PU	random_open_loop	OR	1.414	1.324	CWRU	noise_aug	IR	0.836	2.846
PU	random_open_loop	IR	0.796	0.624	CWRU	noise_aug	B	0.877	4.967
PU	noise_aug	healthy	0.371	—	CWRU	noise_aug	OR	1.516	3.555
PU	noise_aug	OR	0.229	3.043	CWRU	random_open_loop	healthy	NA	—
PU	noise_aug	IR	0.221	2.045	CWRU	random_open_loop	IR	NA	NA
CWRU	llm	healthy	0.150	—	CWRU	random_open_loop	B	NA	NA
CWRU	llm	IR	5.354	2.256	CWRU	random_open_loop	OR	NA	NA

Table 9

Per-class synthetic-to-real audit for the exact physics-v3 pool selections. NRMSE uses class-reference per-channel standardization; $|\rho|_{\max}$ is global-energy-normalized multichannel cross-correlation maximized over all linear lags. Exact denotes byte-identical float32 windows. No distance or correlation threshold is used to infer copying. Berkeley is exemplar-background simulation, so its similarity values characterize retained carrier structure rather than independent waveform generation.

dataset	pool	class	synthetic/reference	exact	min NRMSE	max $ \rho $
PU	llm	healthy	150/1200	0	0.405	0.998
	llm	OR	150/1202	0	0.348	0.998
	llm	IR	150/1444	0	0.657	0.997
	rule	healthy	150/1200	0	0.500	0.997
	rule	OR	150/1202	0	0.440	0.997
	rule	IR	150/1444	0	0.744	0.997
	random_open_loop	healthy	150/1200	0	0.428	0.996
	random_open_loop	OR	150/1202	0	0.353	0.992
	random_open_loop	IR	150/1444	0	0.574	0.997
	noise_aug	healthy	150/1200	0	0.032	1.000
	noise_aug	OR	150/1202	0	0.028	1.000
	noise_aug	IR	150/1444	0	0.032	1.000
CWRU	llm	healthy	20/83	0	1.284	0.300
	llm	IR	20/163	0	0.650	0.281
	llm	B	20/163	0	0.149	0.252
	llm	OR	20/287	0	0.788	0.301
	rule	healthy	20/83	0	0.925	0.671
	rule	IR	20/163	0	0.610	0.320
	rule	B	20/163	0	0.142	0.240
	rule	OR	20/287	0	0.775	0.347
	random_open_loop	all	—	—	—	—
	noise_aug	healthy	20/83	0	0.260	0.972
	noise_aug	IR	20/163	0	0.037	1.000
	noise_aug	B	20/163	0	0.018	1.000
	noise_aug	OR	20/287	0	0.038	1.000
Berkeley	llm	healthy	20/629	0	0.310	0.932
	llm	degraded	20/1207	0	0.425	0.989
	rule	healthy	20/629	0	0.333	0.969
	rule	degraded	20/1207	0	0.337	0.999
	random_open_loop	healthy	20/629	0	3.925	0.826
	random_open_loop	degraded	20/1207	0	1.494	0.806
	noise_aug	healthy	20/629	0	0.043	0.999
	noise_aug	degraded	20/1207	0	0.035	0.999

performance remained unsuccessful. v3 stopped when train-condition morphology mapping could not form an admissible balanced pool. v4 stopped after both train-only candidate-pool strategies failed admission, and v5 showed that its candidate criterion also admitted wrong-label/noise controls. v6 stopped because source-only cyclostationary coherence failed to separate all four held-out targets. The formal held-out evaluations are v1 and v2; v3–v6 remain development or admission stops.




Figure 7: PU closed-loop admission across feedback rounds. (a) Cumulative percentage of proposal slots admitted by or before each round, shown for healthy, outer-race, and inner-race classes ($n = 150$ slots each) and for all slots ($n = 450$); labels give pooled cumulative counts. (b) Newly admitted slots assigned to their first passing round and stacked by class. A passing round is the earliest archived candidate with `feasible=true`; the aggregate exactly reproduces the frozen 286 admitted slots, leaving 164 unadmitted at $K = 3$. Counts describe one frozen pool-generation run and have no independent-pool sampling interval. Source data are provided with the figure artifacts.

Table 10

PU verifier-gate ablation from the cached 450-slot/class $K = 3$ archive. Each evaluable condition uses the frozen file split, 150 synthetic windows/class, and 20 paired CNN seeds. Entries are mean downstream deltas (Accuracy/Macro-F1) from the full gate; PSD-W1 is class-averaged. $2.0\delta_y$ cannot form the fixed healthy-class budget and is not evaluated at a reduced budget.

setting	slot accept.	PSD-W1	$n = 5$ $\Delta A/\Delta F1$	$n = 10$ $\Delta A/\Delta F1$	$n = 25$ $\Delta A/\Delta F1$	status
full gate	0.636	236.9	+0.000/+0.000	+0.000/+0.000	+0.000/+0.000	reference
without M2	0.824	242.1	-0.005/-0.010	-0.012/-0.014	-0.003/-0.004	B=150/class
without M3	0.684	263.5	-0.014/-0.014	-0.012/-0.013	-0.001/-0.002	B=150/class
without M4	0.744	225.0	-0.016/-0.018	-0.012/-0.014	-0.010/-0.011	B=150/class
without M5	0.636	228.0	-0.019/-0.021	-0.013/-0.015	-0.009/-0.010	B=150/class
$0.5\delta_y$	0.636	228.0	-0.019/-0.021	-0.013/-0.015	-0.009/-0.010	same as no M5
$1.0\delta_y$	0.636	236.9	+0.000/+0.000	+0.000/+0.000	+0.000/+0.000	same pool
$2.0\delta_y$	0.636	—	—	—	—	healthy 64/150

6.2. Incomplete protocol metadata stops broader milling claims

The core milling evidence is limited to datasets with complete, auditable protocol metadata. The private machine-tool dataset used in earlier drafts therefore leaves the core evidence. Its speed, geometry, file-to-operation manifest, and audited physical label mapping are incomplete, which prevents a public reproducible portability claim.

UMich and MU-TCM define two distinct protocol stops. In the available UMich protocol, experimental condition and process metadata are confounded with wear labels, allowing a classifier to exploit condition identity. The current files leave these effects entangled, so a valid formal test requires repeated conditions or a revised acquisition design. MU-TCM v3 stops at train-only inner validation, where only 2/6 predeclared comparisons pass; preregistration and formal held-out evaluation were therefore withheld. Both datasets remain boundary evidence.




Figure 8: Non-exclusive PU gate-failure rates by class and recipe source. (a) Cached LLM $K = 3$ rescreen. (b) Random recipes passed through the same verifier. Statistics, soft-spectrum, PSD-W1, and multi-band envelope failures are computed from frozen slot reports.

7. Discussion

7.1. Generation, admission, and diagnosis provide complementary evidence

BREEZE generates synthetic fault waveforms without fitting a target-specific generator. Prompt-supported LLM recipes externalize the controllable signal description, the fixed seeded renderer makes waveform construction reproducible, and the verifier closes the generation loop through train-calibrated admission, rejection, and bounded revision. The resulting admitted pools, rather than the verifier alone, are the framework output tested for physical fidelity and downstream diagnostic utility.

The source ablations answer the first question by establishing an LLM recipe contribution within the controlled BREEZE pipeline. On PU, random recipes subjected to the verifier produce 0 admitted slots and therefore constitute a capacity failure: no balanced downstream pool exists for an effectiveness comparison. The distinct random open-loop control is substantially weaker downstream. LLM recipes also exceed the fixed rule source in every registered PU cell and every provenance-valid CWRU family.

The admission records answer the second question by showing how the verifier shapes feasible pools. Random recipes subjected to the verifier form no balanced pool, whereas accepted LLM pools exhibit waveform, envelope, PSD, band-energy, statistical-distance, and synthetic-to-real similarity records under the frozen diagnostics. Measured violations also drive the next bounded feedback round.

The downstream protocols answer the third question by locating both stable utility and its transfer boundary. LLM recipes improve the registered PU and provenance-valid CWRU comparisons and provide a qualified lower-shot Berkeley advantage. Their extra value converges with the rule recipe at ten Berkeley shots, while PU LOCO remains unsuccessful across six predeclared stages.

Together, these answers support the proposed knowledge-residence change: structured recipes and fixed equations propose signals, while calibrated evidence controls admission. The supported operating point requires few-shot real calibration and provides inspectable records without training a new target generator.

7.2. Knowledge residence guides the choice of generator

The available target evidence determines when the BREEZE operating point is appropriate. By removing target-specific generator fitting and checkpoint selection from augmentation, it fits deployments with few labelled target faults, unavailable or undesirable dedicated generator-training infrastructure, and a requirement for sample-level traceability. The real outer-training split remains necessary for gate calibration, and the downstream diagnostic model still requires training.




Figure 9: Cross-condition evidence. (a,c) Signed CWRU LLM-rule Accuracy and Macro-F1 effects in percentage points when source-load0 pools are evaluated on held-out loads 1–3, on one zero-centred scale; filled dots mark registered Holm passes. The archived held-out-load0 fold is excluded for target-load provenance. (b,d) Number of registered comparisons out of four passing Holm with positive effect in each condition-shot-metric cell for PU LOCO v1 and v2, on a separate five-state discrete scale.

A trained generator becomes a reasonable alternative when a representative target corpus and compute budget support distribution fitting, hyperparameter selection, and checkpoint selection, especially when important morphology cannot be expressed by the fixed renderer. The present evidence neither locates this crossover nor supports a benchmark ranking; the guidance follows the two paradigms’ operational requirements.

7.3. Physical admission evidence versus physical truth

Train-calibrated gates provide auditable physical evidence within a finite set of observables. A candidate can satisfy RMS, PSD, and envelope gates while missing load-path, transfer-function, non-stationary, or degradation mechanisms; a rare real sample can likewise fall outside a finite training boundary. BREEZE therefore reports class-wise distances, downstream tests, and negative transfer results alongside gate pass rates. A pass means “admissible under the declared training-calibrated gates,” not proof of membership in the full real distribution.

The rule comparison partially addresses the circularity created by shared renderer and verifier kinematics because rule and LLM recipes encounter the same renderer, physical checks, and downstream protocol. The random-plus-verifier audit also encounters those checks but stops at a zero-admission capacity failure; the separate random open-loop pool tests non-semantic recipe values without claiming verifier matching. These controls show that shared frequency definitions alone do not explain the useful LLM pool, but they cannot remove the circularity. External expert scoring, observables absent from the prompt, and new condition-disjoint datasets are required for a stronger physical-validity claim.


PU leave-one-condition-out: the complete predeclared failure chain

Figure 10: Complete PU LOCO failure chain. v1 and v2 are formal held-out evaluations; v3–v6 stop at train-only morphology, admission, sanity, or source-only evidence gates. The sequence prevents unsuccessful development stages from being reported as completed cross-condition tests.

7.4. Reproducibility and remaining blockers

The repository supports replay from archived recipes through reported statistics. It stores split manifests, prompts, recipe JSON, local renderer seeds, waveform hashes, per-gate reports, pool manifests, seed-level classifier rows, registered tests, and table-generation assertions. Long-running pipelines are append-only or checkpoint-resumable. Frozen CSV/NPZ artifacts generate every reported number, and NA explicitly marks an unavailable pool.

Two provenance and baseline blockers remain. The original PU archive omits the provider version and provider-side sampling seed, so exact replay begins from the archived recipes instead of the original LLM response. The present study does not establish a performance or cost ranking against fully trained generative models because matched formal TimeGAN/DDPM comparisons are unavailable. Additional diagnostic-backbone comparisons are also unavailable, which keeps universal classifier claims outside the supported scope. Future work should complete those matched baselines, repeat generation to support independent-pool inference, test at least one non-CNN diagnostic backbone, evaluate cross-specimen and cross-machine splits, and preregister a new PU cross-condition representation before accessing held-out outcomes. The present seed-level tests quantify repeated few-shot/CNN variability around one shared fixed pool per method and do not estimate between-pool generation variability.

8. Conclusion

BREEZE demonstrates that physically structured and diagnostically useful synthetic fault waveforms can be generated without fitting a target-specific generator. Prompt-supported LLM recipes, a deterministic seeded renderer, and a train-calibrated verifier convert augmentation into an explicit and auditable generation process with bounded revision and traceable records.

Across three frozen public few-shot protocols, this design isolates an LLM recipe contribution on PU, improves the registered provenance-valid CWRU comparisons, and provides a qualified lower-shot Berkeley benefit. The physical audit further shows that the PU LLM pool improves RMS, PSD, and band-energy fidelity relative to rule and random recipe sources while retaining comparator-favouring kurtosis and alignment evidence. Six predeclared PU cross-condition attempts delimit the current morphology-transfer boundary.

The resulting contribution is a practical operating point for training-free augmentation in which physical controls, generation decisions, and diagnostic utility remain inspectable. Future work should complete matched trained-generator comparisons, repeat independent pool generation, and evaluate additional diagnostic backbones and cross-machine protocols.

Data and code availability

PU and CWRU are public benchmarks, and Berkeley milling is available through the NASA prognostics repository subject to its terms. The Git repository contains the prompt builder and schema, deterministic renderer, verifier, table/figure scripts, frozen seed-level statistical outputs, PU pool arrays and manifests, physics manifests with SHA-256 hashes, and the evidence ledger. Primary CWRU and Berkeley waveform pools, recipes, and detailed candidate records currently reside in ignored `breeze/runs/` paths and are not part of the versioned Git tree. They will be deposited as a checksum-indexed repository release or data-repository supplement for exact waveform replay; until that deposit, the Git repository supports numerical-result audit and PU pool replay, but not complete CWRU/Berkeley waveform replay. API credentials, provider-side sampling state, and raw third-party datasets remain outside the release.

Declaration of competing interest

The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Declaration of generative AI and AI-assisted technologies

During preparation of this work, the authors used OpenAI Codex for code inspection, experiment orchestration, figure and table assembly, and language editing. The authors reviewed the evidence, verified the reported results, and take responsibility for the manuscript.

References

- Agogino, A., Goebel, K., 2007. Milling data set. NASA Ames Prognostics Data Repository. URL: <https://www.nasa.gov/content/prognostics-center-of-excellence-data-set-repository>.
- Alcaraz, J.M.L., Strodthoff, N., 2022. Diffusion-based time series imputation and forecasting with structured state space models. arXiv preprint arXiv:2208.09399.
- Antoni, J., 2006. The spectral kurtosis: A useful tool for characterising non-stationary signals. *Mechanical Systems and Signal Processing* 20, 282–307. doi:10.1016/j.ymssp.2004.09.001.
- Antoni, J., 2007. Fast computation of the kurtogram for the detection of transient faults. *Mechanical Systems and Signal Processing* 21, 108–124. doi:10.1016/j.ymssp.2005.12.002.
- Benjamini, Y., Hochberg, Y., 1995. Controlling the false discovery rate: A practical and powerful approach to multiple testing. *Journal of the Royal Statistical Society: Series B* 57, 289–300.
- Desai, A., Freeman, C., Wang, Z., Beaver, I., 2021. TimeVAE: A variational auto-encoder for multivariate time series generation. arXiv preprint arXiv:2111.08095.
- Esteban, C., Hyland, S.L., Rätsch, G., 2017. Real-valued (medical) time series generation with recurrent conditional GANs. arXiv preprint arXiv:1706.02633.
- Feng, Y., Chen, J., Xie, J., Zhang, T., Lv, H., Pan, T., 2022. Meta-learning as a promising approach for few-shot cross-domain fault diagnosis: Algorithms, applications, and prospects. *Knowledge-Based Systems* 235, 107646. doi:10.1016/j.knosys.2021.107646.
- Gao, Q., Xu, T., Zhang, S., Fang, Y., Tan, H., 2025. Bearing fault signal generation method by fusion of physical constraints and multimodal features. *Measurement Science and Technology* 36, 116109. doi:10.1088/1361-6501/ae1a04.
- Gretton, A., Borgwardt, K.M., Rasch, M.J., Schölkopf, B., Smola, A., 2012. A kernel two-sample test. *Journal of Machine Learning Research* 13, 723–773.
- Gruver, N., Finzi, M., Qiu, S., Wilson, A.G., 2023. Large language models are zero-shot time series forecasters. *Advances in Neural Information Processing Systems* 36, 19622–19635.
- Guo, Z., Xu, L., Zheng, Y., Xie, J., Wang, T., 2025. Bearing fault diagnostic framework under unknown working conditions based on condition-guided diffusion model. *Measurement* 242, 115951. doi:10.1016/j.measurement.2024.115951.
- Holm, S., 1979. A simple sequentially rejective multiple test procedure. *Scandinavian Journal of Statistics* 6, 65–70.
- Kapoor, S., Narayanan, A., 2023. Leakage and the reproducibility crisis in machine-learning-based science. *Patterns* 4, 100804. doi:10.1016/j.patter.2023.100804.
- Kim, T., Ko, J.U., Lee, J., Kim, Y.C., Jung, J.H., Youn, B.D., 2024. Spectrum-guided GAN with density-directionality sampling: Diverse high-fidelity signal generation for fault diagnosis of rotating machinery. *Advanced Engineering Informatics* 62, 102821. doi:10.1016/j.aei.2024.102821.
- Ko, J.U., Lee, J., Kim, T., Kim, Y.C., Youn, B.D., 2023. Frequency-learning generative network (FLGN) to generate vibration signals of variable lengths. *Expert Systems with Applications* 227, 120255. doi:10.1016/j.eswa.2023.120255.
- Lee, J., Jeon, H., Kim, U., Kim, M., 2026. BearGen: LLM-guided signal generation framework for bearing fault diagnosis. *Advanced Engineering Informatics* 71, 104400. doi:10.1016/j.aei.2026.104400.
- Lessmeier, C., Kimotho, J.K., Zimmer, D., Sestro, W., 2016. Condition monitoring of bearing damage in electromechanical drive systems by using motor current signals of electric motors: A benchmark data set for data-driven classification. *PHM Society European Conference* 3. doi:10.36001/phme.2016.v3i1.1577.

- Li, C., Li, S., Zhang, A., He, Q., Liao, Z., Hu, J., 2021. Meta-learning for few-shot bearing fault diagnosis under complex working conditions. *Neurocomputing* 439, 197–211. doi:10.1016/j.neucom.2021.01.099.
- Li, X., Metsis, V., Wang, H., Ngu, A.H.H., 2022a. TTS-GAN: A transformer-based time-series generative adversarial network, in: *Artificial Intelligence in Medicine*, Springer. pp. 133–143. doi:10.1007/978-3-031-09342-5_13.
- Li, X., Ngu, A.H.H., Metsis, V., 2022b. TTS-CGAN: A transformer time-series conditional GAN for biosignal data augmentation. *arXiv preprint arXiv:2206.13676*.
- Liao, S., Ni, H., Szpruch, L., Wiese, M., Sabate-Vidales, M., Xiao, B., 2020. Conditional Sig-Wasserstein GANs for time series generation. *arXiv preprint arXiv:2006.05421*.
- Lin, C., Kong, Y., Zhao, K., Han, Q., Dong, M., Liu, H., Chu, F., 2025. Filling generative adversarial network: A novel intelligent machinery diagnostic method towards extremely limited data. *Advanced Engineering Informatics* 68, 103666. doi:10.1016/j.aei.2025.103666.
- Liu, F., Han, Q., Chen, J., Ding, B., Zhang, G., Shen, Z., Wu, X., Zhang, X., Li, W., 2026. Physics-informed diffusion-based augmentation for vibration-based fault diagnosis of rotating machinery. *Journal of Vibration and Control* doi:10.1177/10775463261455715.
- Liu, S., Jiang, H., Wu, Z., Li, X., 2022. Data synthesis using deep feature enhanced generative adversarial networks for rolling bearing imbalanced fault diagnosis. *Mechanical Systems and Signal Processing* 163, 108139. doi:10.1016/j.ymssp.2021.108139.
- Ma, L., Ding, Y., Wang, Z., Wang, C., Ma, J., Lu, C., 2021. An interpretable data augmentation scheme for machine fault diagnosis based on a sparsity-constrained generative adversarial network. *Expert Systems with Applications* 182, 115234. doi:10.1016/j.eswa.2021.115234.
- Nichol, A.Q., Dhariwal, P., 2021. Improved denoising diffusion probabilistic models, in: *International Conference on Machine Learning*, PMLR. pp. 8162–8171.
- Peng, H., Liu, J., Du, J., Gao, J., Wang, W., 2025. BearLLM: A prior knowledge-enhanced bearing health management framework with unified vibration signal representation, in: *Proceedings of the AAAI Conference on Artificial Intelligence*, pp. 19866–19874. doi:10.1609/aaai.v39i19.34188.
- Raissi, M., Perdikaris, P., Karniadakis, G.E., 2019. Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations. *Journal of Computational Physics* 378, 686–707. doi:10.1016/j.jcp.2018.10.045.
- Randall, R.B., Antoni, J., 2011. Rolling element bearing diagnostics—a tutorial. *Mechanical Systems and Signal Processing* 25, 485–520. doi:10.1016/j.ymssp.2010.07.017.
- Rasul, K., Seward, C., Schuster, I., Vollgraf, R., 2021. Autoregressive denoising diffusion models for multivariate probabilistic time series forecasting, in: *International Conference on Machine Learning*, PMLR. pp. 8857–8868.
- Schoen, R.R., Habetler, T.G., Kamran, F., Bartheld, R.G., 1995. Motor bearing damage detection using stator current monitoring. *IEEE Transactions on Industry Applications* 31, 1274–1279. doi:10.1109/28.475697.
- Smith, W.A., Randall, R.B., 2015. Rolling element bearing diagnostics using the Case Western Reserve University data: A benchmark study. *Mechanical Systems and Signal Processing* 64–65, 100–131. doi:10.1016/j.ymssp.2015.04.021.
- Su, H., Xiang, L., Hu, A., Xu, Y., Yang, X., 2022. A novel method based on meta-learning for bearing fault diagnosis with small sample learning under different working conditions. *Mechanical Systems and Signal Processing* 169, 108765. doi:10.1016/j.ymssp.2021.108765.
- Tashiro, Y., Song, J., Song, Y., Ermon, S., 2021. CSDI: Conditional score-based diffusion models for probabilistic time series imputation, in: *Advances in Neural Information Processing Systems*, pp. 24804–24816.
- Tong, Q., Zhu, R., Lu, F., Wang, B., Du, S., Jiang, X., Xu, J., Huo, J., Du, X., 2026. ACS-DM: Adaptive conditional sampling diffusion model for few-shot machinery fault diagnosis. *Journal of Vibration and Control* doi:10.1177/10775463251414180.
- Villani, C., 2009. *Optimal Transport: Old and New*. Springer, Berlin. doi:10.1007/978-3-540-71050-9.
- Welch, P.D., 1967. The use of fast Fourier transform for the estimation of power spectra: A method based on time averaging over short, modified periodograms. *IEEE Transactions on Audio and Electroacoustics* 15, 70–73. doi:10.1109/TAU.1967.1161901.
- Wheat, L., Mohrenschildt, M.V., Habibi, S., Al-Ani, D., 2024. Impact of data leakage in vibration signals used for bearing fault diagnosis. *IEEE Access* 12, 169879–169895. doi:10.1109/ACCESS.2024.3497716.
- Wilcoxon, F., 1945. Individual comparisons by ranking methods. *Biometrics Bulletin* 1, 80–83. doi:10.2307/3001968.
- Wu, J., Zhao, Z., Sun, C., Yan, R., Chen, X., 2020. Few-shot transfer learning for intelligent fault diagnosis of machine. *Measurement* 166, 108202. doi:10.1016/j.measurement.2020.108202.
- Xu, J., Tong, Q., Zhu, R., Du, S., Zhao, J., Jiang, X., Wang, B., 2026. A diffusion-based data augmentation framework for few-shot fault diagnosis of intelligent high-speed train components. *Sensors* 26, 3091. doi:10.3390/s26103091.
- Xu, Y., Chen, Z., Wang, R., Li, Y., Tang, F., Zhao, M., Liu, J., 2025. FaultDiffusion: Few-shot fault time series generation with diffusion model. *arXiv preprint arXiv:2511.15174* doi:10.48550/arXiv.2511.15174.
- Yang, X., Ye, T., Yuan, X., Zhu, W., Mei, X., Zhou, F., 2024. A novel data augmentation method based on denoising diffusion probabilistic model for fault diagnosis under imbalanced data. *IEEE Transactions on Industrial Informatics* 20, 7820–7831. doi:10.1109/TII.2024.3366991.
- Yi, H., Hou, L., Jin, Y., Saeed, N.A., Kandil, A., Duan, H., 2024. Time series diffusion method: A denoising diffusion probabilistic model for vibration signal generation. *Mechanical Systems and Signal Processing* 216, 111481. doi:10.1016/j.ymssp.2024.111481.
- Yoon, J., Jarrett, D., van der Schaar, M., 2019. Time-series generative adversarial networks, in: *Advances in Neural Information Processing Systems*, pp. 5508–5518.
- Yu, T., Li, C., Huang, J., Xiao, X., Zhang, X., Li, Y., Fu, B., 2024. ReF-DDPM: A novel DDPM-based data augmentation method for imbalanced rolling bearing fault diagnosis. *Reliability Engineering & System Safety* 251, 110343. doi:10.1016/j.res.s.2024.110343.